Overall Paired Differences (Incremental - Comparator)

Method	Metric	Difference	95% CI	p	Holm p
vs. Waterfall	E2E (pp)	+51.1	[+41.1, +60.0]	<0.001	<0.001
vs. Waterfall	PFS (points)	+1.80	[+1.43, +2.17]	<0.001	<0.001
vs. Waterfall	MFS@4 (pp)	+34.3	[+25.6, +43.6]	<0.001	<0.001
vs. Waterfall	Tokens	-126,633.5	[-143,764.6, -108,943.8]	<0.001	<0.001
vs. Incremental w/o Self-Refinement	E2E (pp)	+62.2	[+54.4, +70.0]	<0.001	<0.001
vs. Incremental w/o Self-Refinement	PFS (points)	+2.20	[+1.80, +2.60]	<0.001	<0.001
vs. Incremental w/o Self-Refinement	MFS@4 (pp)	+43.5	[+34.3, +52.6]	<0.001	<0.001
vs. Incremental w/o Self-Refinement	Tokens	+52,017.6	[+43,228.1, +60,882.5]	<0.001	<0.001

Incremental minus comparator; n=30 paired tasks x 3 selected runs. Positive quality / negative Token effects favor Incremental; pp denotes percentage points.
95% task-cluster bootstrap CIs (30,000 samples) condition on observed repeats. Two-sided task-level permutation tests: 199,999 permutations, plus-one correction.
Holm p adjusts all 8 overall comparisons; CIs are pointwise, not Holm-adjusted. * Zero-width CIs do not imply certainty. Full methodological notes accompany the MD/TeX exports.