## Overall Estimates, Run Variability, and Confidence Intervals

| Method | Metric | Estimate | Run SD | 95% CI |
|---|---|---|---|---|
| **Incremental** | E2E | 92.2% | 6.9 pp | [86.7%, 96.7%] |
| **Incremental** | PFS (points) | 3.33 | 0.28 | [3.00, 3.65] |
| **Incremental** | MFS@4 | 52.1% | 3.3 pp | [43.6%, 60.3%] |
| **Incremental** | Tokens | 161,958.7 | 27,608.4 | [151,115.6, 173,131.5] |
| **Waterfall** | E2E | 41.1% | 6.9 pp | [32.2%, 50.0%] |
| **Waterfall** | PFS (points) | 1.53 | 0.24 | [1.19, 1.89] |
| **Waterfall** | MFS@4 | 17.8% | 2.7 pp | [11.6%, 24.2%] |
| **Waterfall** | Tokens | 288,592.1 | 33,852.4 | [270,455.1, 305,782.9] |
| **Incremental w/o Self-Refinement** | E2E | 30.0% | 8.8 pp | [22.2%, 37.8%] |
| **Incremental w/o Self-Refinement** | PFS (points) | 1.13 | 0.45 | [0.82, 1.45] |
| **Incremental w/o Self-Refinement** | MFS@4 | 8.6% | 5.3 pp | [4.9%, 12.8%] |
| **Incremental w/o Self-Refinement** | Tokens | 109,941.0 | 12,728.6 | [105,041.5, 114,974.6] |

n=30 paired tasks x 3 selected runs per method. Run SD is descriptive sample SD across run aggregates (ddof=1), not SE; rate SDs use pp.
Pointwise 95% CIs: 30,000 paired, complexity-stratified task-cluster bootstrap samples, conditional on observed repeats. MFS@4 pools modules.
* Zero-width empirical CIs do not imply certainty. E2E covers build/compilation and startup, not complete business functionality.
Full statistical notes accompany the MD/TeX exports.